

SmolTLV: A Small TLV Binary Encoding

Contents

1.	Status of This Document	2
2.	Overview	2
3.	Goals	2
4.	Data Model	2
5.	Wire Format	3
5.1.	Item Encoding	3
5.2.	Length Encoding	3
5.3.	Endianness	3
6.	Integer Encoding	3
7.	Bytes and Strings	3
7.1.	Bytes	3
7.2.	String	3
8.	Containers	3
8.1.	List	3
8.2.	Dict	4
8.2.1.	Key Type	4
8.2.2.	Duplicate Keys	4
9.	Canonical / Deterministic Encoding (Optional Profile)	4
10.	Error Handling	4
11.	Security Considerations	4
12.	Example Encodings	5
12.1.	A null value	5
12.2.	Boolean values	5
12.3.	Integers	5
12.4.	Simple string	5
12.5.	List	5
12.6.	Dictionary	6
13.	Appendix: Type Registry	6

1. Status of This Document

This document is a draft specification for **SmolTLV**, a compact TLV (Type–Length–Value) binary encoding designed for simple incremental parsing on resource-constrained systems.

Implementations **MUST** follow the normative requirements described using capitalized keywords (**MUST**, **SHOULD**, **MAY**), interpreted as in RFC 2119.

2. Overview

SmolTLV encodes a sequence of **items**. Each item is self-delimiting and consists of:

- a 1-byte **Type**,
- a 24-bit big-endian **Length** (payload length in bytes),
- a **Payload** of exactly **Length** bytes.

The format supports both primitive values (e.g. integers, byte strings) and containers (lists and dictionaries) whose payload is a concatenation of nested SmolTLV items.

3. Goals

- **Incremental parsing**: Decoders can parse items from a streaming source and skip unknown types.
- **Memory mapping friendly**: Length-delimited items allow walking a buffer without copying.
- **Small and obviously-correct implementations**: Fixed-size header (4 bytes) and no indefinite-length constructs.

4. Data Model

SmolTLV defines the following built-in types:


Type name	Value	Meaning
Null	0x00	An explicit null value. Payload length MUST be 0.
True	0x01	Boolean true. Payload length MUST be 0.
False	0x02	Boolean false. Payload length MUST be 0.
Int	0x03	Signed integer. Payload length MUST be 8.
Bytes	0x04	Byte string (opaque blob).
String	0x05	UTF-8 text string.
List	0x06	Ordered sequence of items (container).
Dict	0x07	Key/value mapping (container).

Type values outside this set are **reserved for extensions**. Decoders **MUST** be able to skip unknown types based solely on the length field.

5. Wire Format

5.1. Item Encoding

Each item has header with following layout (in big-endian byte order):



Where:

- **Type** is an unsigned 8-bit value.
- **Length** is an unsigned 24-bit integer, encoded big-endian, giving the number of payload bytes.
- Maximum payload length is 0xFFFFFFF (16,777,215) bytes.

Followed by exactly **Length** bytes of payload data.

5.2. Length Encoding

Length is encoded as three bytes:

```
len = (b1 << 16) | (b2 << 8) | b3
```

where b1, b2, b3 are the three length bytes in network (big-endian) order.

5.3. Endianness

All multi-byte integers in SmolTLV payloads **MUST** be encoded big-endian unless otherwise specified.

6. Integer Encoding

The `SMOLTLV_TYPE_INT` payload encodes a signed 64 bit integer using **two's complement big-endian**.

Payload length must be 8 bytes.

7. Bytes and Strings

7.1. Bytes

`SMOLTLV_TYPE_BYTES` payload is an opaque byte string of length `Length`. No interpretation is implied.

7.2. String

`SMOLTLV_TYPE_STRING` payload **MUST** be valid UTF-8. Implementations **MAY** reject invalid UTF-8.

Canonical form (if enforced by a profile) **SHOULD** use Unicode NFC, but SmolTLV does not require normalization by default.

8. Containers

8.1. List

`SMOLTLV_TYPE_LIST` payload is a concatenation of zero or more complete SmolTLV items.

A decoder processes a list by creating a bounded cursor over the list payload and repeatedly parsing items until the cursor reaches the end of the payload.

The list payload **MUST NOT** contain partial items or padding; i.e., the nested items **MUST** exactly fill the list payload length.

8.2. Dict

SMOLTLV_TYPE_DICT payload is a concatenation of SmolTLV items representing key/value pairs:

- The payload **MUST** contain an even number of nested items.
- Items are interpreted as (key1, value1, key2, value2, ...).

8.2.1. Key Type

By default, SmolTLV permits any item type as a dictionary key. However, many deployments **SHOULD** restrict keys to SMOLTLV_TYPE_STRING for interoperability.

If a deployment profile restricts dict keys, encoders **MUST** follow that profile and decoders **SHOULD** enforce it.

8.2.2. Duplicate Keys

SmolTLV does not define duplicate-key semantics. A profile **MUST** specify one of:

- reject duplicates,
- last-wins,
- first-wins,
- collect-all.

9. Canonical / Deterministic Encoding (Optional Profile)

SmolTLV itself allows multiple equivalent encodings (e.g., dictionary order, integer minimality is required but key ordering is not). For deterministic encoding, a profile **MAY** require:

- Dict keys **MUST** be strings.
- Dict entries **MUST** be sorted lexicographically by the UTF-8 byte sequence of the key.
- Duplicate keys **MUST** be rejected.
- Strings **MUST** be valid UTF-8.

If deterministic encoding is required, all producers and consumers **MUST** implement the same profile.

10. Error Handling

Decoders **SHOULD** distinguish:

- **Need more data:** insufficient bytes to complete the header or payload.
- **Format error:** invalid structure (e.g., container payload not fully consumable into items, invalid minimal integer encoding, invalid UTF-8 if enforced).

Decoders **MUST** be able to skip unknown types by advancing the cursor by 4 + Length bytes.

11. Security Considerations

Implementations **MUST** defend against:

- Length fields that exceed available buffer space (bounds checks).
- Excessive nesting depth (use a maximum depth).
- Excessive total allocation (cap sizes when materializing structures).
- Quadratic behavior when processing dictionaries with adversarial key patterns.

12. Example Encodings

12.1. A null value

The null value:

- Type = Null (0x00)
- Length = 0
- Payload = (none)

00 00 00 00

12.2. Boolean values

The boolean true value:

- Type = Bool True (0x01)
- Length = 0
- Payload = (none)

01 00 00 00

12.3. Integers

The integer value 42:

- Type = Int (0x03)
- Length = 8
- Payload = 0x00 0x00 0x00 0x00 0x00 0x00 0x00 0x2A

03 00 00 08 00 00 00 00 00 00 00 2A

The integer value -1:

- Type = Int (0x03)
- Length = 8
- Payload = 0xFF 0xFF 0xFF 0xFF 0xFF 0xFF 0xFF 0xFF

03 00 00 08 FF FF FF FF FF FF FF FF

12.4. Simple string

The value "hi":

- Type = String (0x05)
- Length = 2
- Payload = 0x68 0x69

05 00 00 02 68 69

12.5. List

The list: ["a", "b"]

06 00 00 0A

05 00 00 01 61

05 00 00 01 62

12.6. Dictionary

The dictionary {"k": 123}

```
07 00 00 0B  
05 00 00 01 6B  
03 00 00 01 7B
```

13. Appendix: Type Registry

The following type assignments are defined by this specification:

```
SMOLTLV_TYPE_NULL      = 0x00,  
SMOLTLV_TYPE_BOOL_TRUE = 0x01,  
SMOLTLV_TYPE_BOOL_FALSE = 0x02,  
SMOLTLV_TYPE_INT       = 0x03,  
SMOLTLV_TYPE_BYTES     = 0x04,  
SMOLTLV_TYPE_STRING    = 0x05,  
SMOLTLV_TYPE_LIST      = 0x06,  
SMOLTLV_TYPE_DICT      = 0x07,
```

Values 0x08..0xFF are reserved for future assignment or application-specific extensions.